



CS 25200: Systems Programming

Lecture 5: File System Hierarchy, Program Structure, Generation, and Loading

Prof. Turkstra



Lecture 05

- UNIX file system structure
- Processes and users
- Program segments
- Executable formats
- Building a program
- Libraries
- Loader

UNIX File System

- *NIX has a hierarchical file system
- Root of entire tree is denoted with /
- Disks and network shares can be **mounted** anywhere inside the hierarchy

Hierarchy

- /bin – historically contained fundamental utilities (ls, cp, etc)
 - Often now a symlink to /usr/bin
- /usr - “user file system”
 - Used to be split off on separate storage, usually not anymore
- /boot – files necessary for startup (e.g., initial kernel image)

- /dev – peripheral and other devices
- /tmp – temporary files that do not survive a reboot
- /var – files that may change frequently (variable)
 - Email, logs, databases, etc
- /sbin - “system (superuser) binaries”
 - Utilities needed to start and maintain/recover the system
 - Also symlinked to /usr/sbin

- /etc – system-wide configuration files
 - Init system scripts
 - /etc/rc.d
 - /etc/rc.*
 - User and group information
 - /etc/passwd, /etc/groups
- /home – user home directories
- /lib – essential libraries
 - Often split into lib, lib64 now
 - Also usually symlinked into /usr/lib*

- /usr/include – header files for libraries and kernel
- /proc – virtual file system that provides information about processes through files
- /sys – other hardware and kernel-related information

Processes

- top, ps (-e, -ax, -f, -u)

Users

- UNIX was designed as a multiuser system
- Database of users exists in /etc/passwd
 - ...but not for Purdue machines (NIS/LDAP)
 - Each user has a unique uid
 - Passwords are now stored in /etc/shadow

- useradd – create a new user
- passwd – change a user’s password
- There exists a special user called “root”
 - Root can do anything
 - Modify files, change permissions/ownership, mount/umount, rm anything, execute anything, etc etc
 - To become root, you use the “su” command. (super user)

Got root?

- Many times systems are configured to use “sudo”. Sudo allows regular users to run certain (specified) things as root. See /etc/sudoers
 - Does not require the root user’s password, unlike su
 - Can “sudo su” to get a root shell, if you have the right privileges
- Root shells are dangerous
 - `$ rm -rf / home/turkstra/somefile`



Groups

- A group is a collection of users
- Users may belong to multiple groups
 - “Supplemental groups”
 - `usermod -S`, `groupadd`, `groupmod`
 - `/etc/group`


Network Information Service (NIS)

- Formerly Yellow Pages or YP
- Client-server directory service protocol for distributing information like users and groups
- See “ypcat”

Program vs. process

- A **program** is an executable file that contains a set of instructions
 - Usually stored on disk or other secondary storage
- A **process** is a program in execution
 - It resides, at least partially, in memory

Process memory layout



32-bit vs 64-bit

- 32-bit systems usually have shared libraries at the lowest address, followed by the text segment
- Starting addresses differ
 - Text or code usually starts 0x400000 on 64-bit, 0x8047000 on 32-bit

Text segment

- Also called the **code** segment
- Contains actual program instructions and any statically linked libraries
- Often **read only** and **executable**
 - Self-modifying code

Demo

■ hello.c

```
#include <unistd.h>
int main(int argc, char *argv[])
{
    int ret = 0xbeefbeef;
    ret = write(1, "Hello\n", 6);
    return 0;
}
```

```
$ gcc -S hello.c
```

```
$ objdump -Dl a.out
```

Data segment

- Initialized global variables and static strings

BSS

- Block started by symbol
- Holds uninitialized global variables
 - By C convention are initialized automatically to 0

Heap

- Dynamically allocated memory
 - i.e., obtained via `malloc()`
- Grows upward as memory is requested
 - Upward → increasing addresses

Stack

- Holds local, temporary, or automatic, variables
- Arguments passed during a function call
- Information needed to return to a previous point in the program
- Grows downward (decreasing addresses)

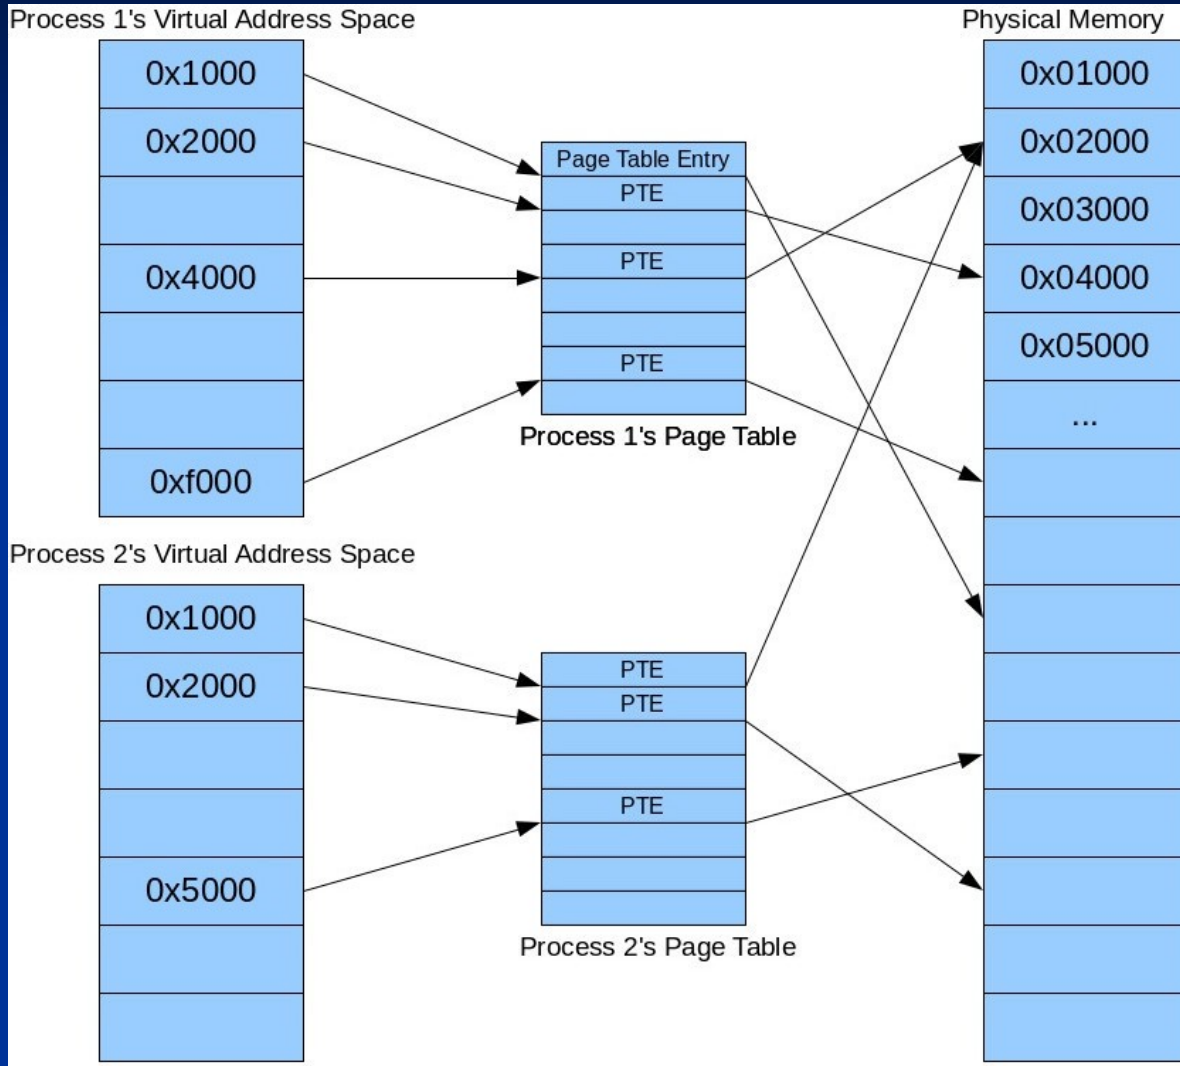
Stack

- A **stack** is a very common data structure used in programs and architectures
- Stacks are generally **LIFO** queues
 - Last in, first out
- Two operations
 - **PUSH** – add something to the stack
 - **POP** – retrieve the most recent item

Shared libraries

- Shared libraries permit the same text segment to be used in multiple processes
 - Through the magic of virtual memory
 - Have their own data and BSS segments
 - Usually copy-on-write

Virtual memory



Demo

- `$ cat /proc/self/maps`

- `howdy.c`:

```
int a = 5;           // Stored in data
int b[20];           // Stored in BSS
int main() {         // Stored in text
    int x;            // Stored in stack
    int *p = (int *) malloc(sizeof(int)); // Heap
    printf("%p\n%p\n%p\n%p\n%p\n", &a, &b, &x, p,
           main);
}
```

Gaps

- Often gaps between segments
- Attempting to access an address in an **unmapped** region causes the OS to send the process a signal: **SIGSEGV**
 - Segmentation violation
 - By default program is immediately terminated and dumps **core**
- This also happens if you access a mapped but **protected** region
 - E.g., try executing in the data segment

Core dump file

- `$ man 5 core`
- File containing an image of the process' memory at time of termination
- Can be used with gdb (or other debuggers)
- May have to enable it (e.g., on `data.cs.purdue.edu`)
 - `$ ulimit -c unlimited`



Program

- File in a particular format containing necessary information to load an application into memory and execute it
 - Often time part of this is split off into the “loader” and libraries
- Programs include:
 - Machine instructions
 - Initialized data
 - List of library dependencies
 - List of memory sections
 - List of values determined at load time

Executable file formats

- Fully “compiled” programs can come in many different formats...
 - ELF – Executable Linux File
 - Used by most recent UNIX systems (e.g., Solaris, Linux)
 - PE – Portable Executable
 - Used by Windoze
 - Mach-O – Mach object [file format]
 - Used by macOS/iOS
 - COFF – Common Object File Format
 - Also windoze, some embedded systems - historically System V Unix
 - a.out – Used in BSD (Berkeley Standard Distribution)
 - Restrictive, rarely seen anymore

ELF

- File header
 - Magic number
 - Version
 - Target ABI
 - ISA
 - Entry point
 - Pointers to
 - Program header
 - Section header
 - etc

Program header

- How to create the process image
 - Segments
 - Types
 - Flags
 - File offset
 - Virtual address
 - Size in file
 - Size in memory

Section header

- Type (data, string, notes, etc)
- Flags (writable, executable, etc)
- Virtual address
- Offset in file image
- Size
- Alignment

- `readelf --headers /bin/ls`
- `objdump -p, -h, -t`

Building a program

- Start with source code
 - abs.c
- Preprocessor
- Compiler
- Assembler
- Linker

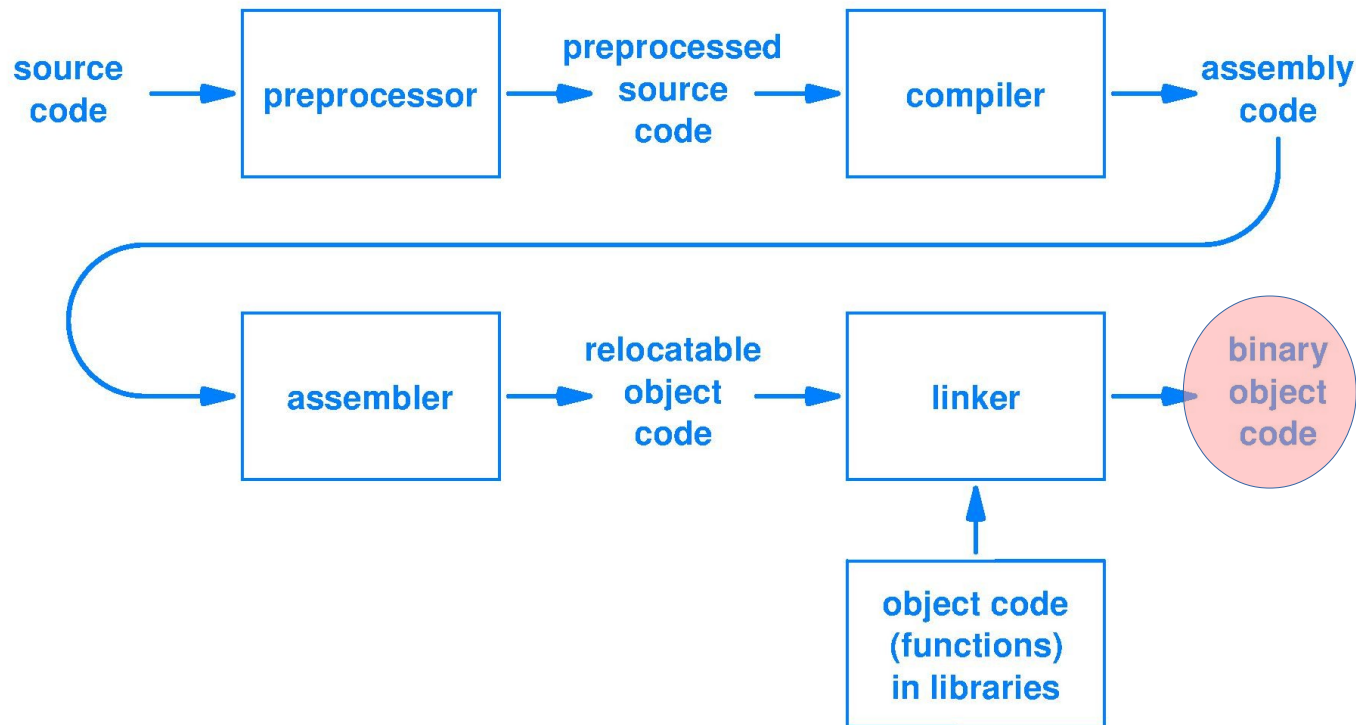


Figure 4.6 The steps used to translate a source program to the binary object code representation used by a processor.

Preprocessor

- When a .c file is compiled, it is first scanned and modified by the preprocessor before being handed to the real compiler
- Finds lines beginning with #, hides them from the compiler, or takes some action
- #include, #define
- #ifdef, #else, #endif

Why macros?

- Run time efficiency
 - No function call overhead
- Passed arguments can be any type
 - `#define MAX(x,y) ((x) > (y) ? (x) : (y))`
 - Works with ints, floats, doubles, even chars

■ Can do math

- `#if (FLAG % 4 == 0) || (FLAG == 13)`

■ Macros

- `#define INC(x) x+1`
- No semi-colon
- Have to be careful
 - `#define ABS(x) x < 0 ? -x : x`
 - `ABS(B+C)`

- Parentheses around substitution variables

```
#define ABS(x) ( (x) < 0 ? -(x) : (x) )
```

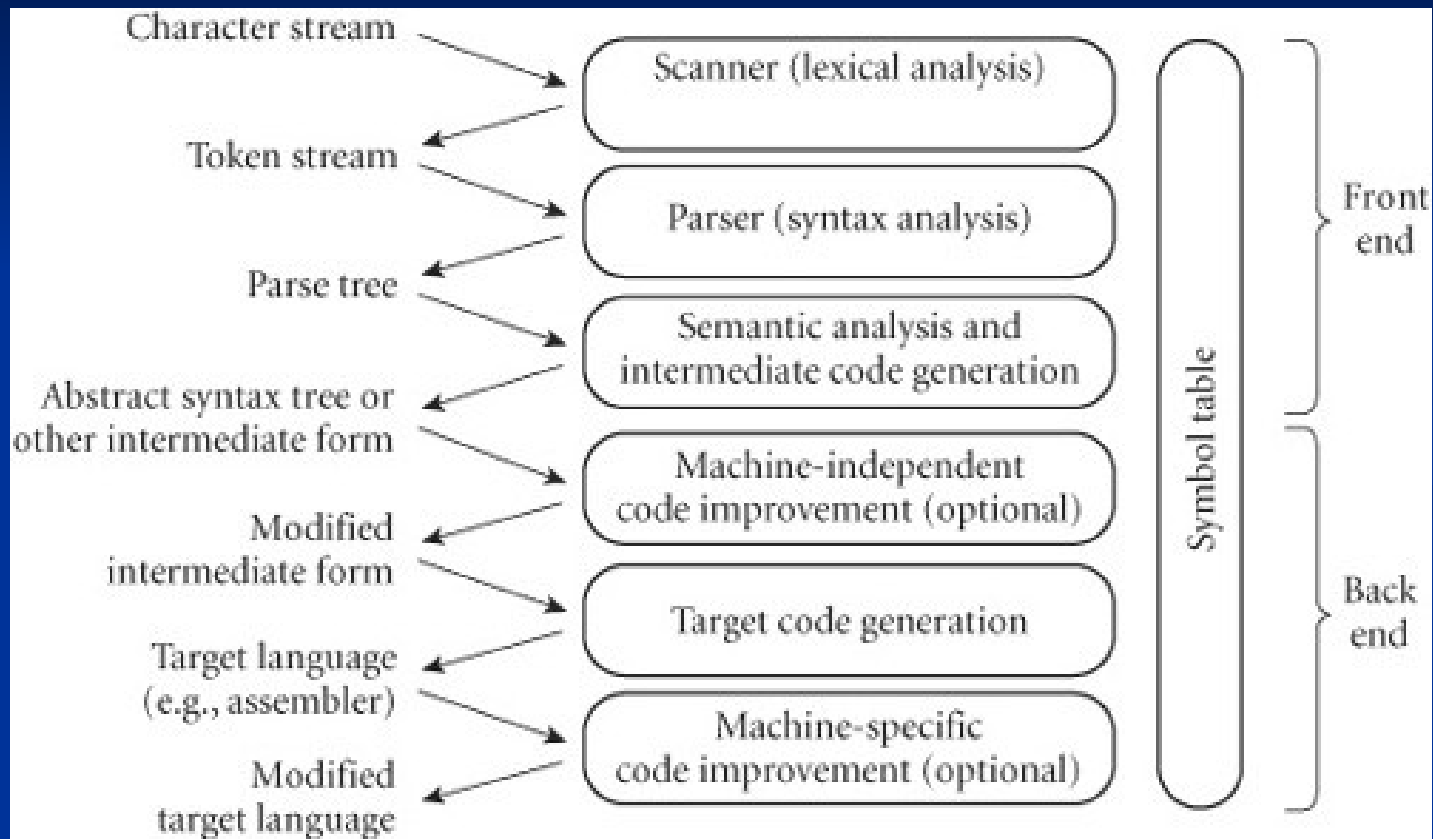

Lots of other tricks

```
printf("The date is %s\n", __DATE__);
```

- Most preprocessor features are used for large/advanced software development practices

- `gcc -E file.c > output.i`

Compiler?



* <http://www.cs.montana.edu/~david.watson5/>

■ gcc -S

Assembler

- Discussed in CS 250 – Computer Architecture
- `gcc -c`
`nm -v`
- Really uses `as`

Libraries

- Libraries are just collections of object files
 - Internal symbols are indexed for fast lookup by the linker
- Searched for symbols that aren't defined in the program
 - Symbol found, pull it into executable (static)
 - Otherwise include a pointer to the file, loaded by loader

Statically linked

- Faster, to a degree
- Portable
- Larger binaries
- Fixed version, no updates
- File extension .a

Dynamically linked

- More complexity
- Easy to upgrade libraries
 - Vulnerabilities
- Have to manage versions
- Loader re-links every time program is executed

```
readelf --dynamic /bin/ls  
ldd /bin/ls
```


End result

- `gcc -o abs abs.c`
`nm -v ./abs`

Loader

- Essential step in starting a program
- Historically allocated space for all sections of the executable (text, data, bss, etc)
- Now simply establishes mappings
 - Page faults actually populate the memory
 - For executable as well as (shared) libraries

- Also resolves any values in the executable to point to the functions/variables in the shared libraries
- Jumps to `_start`
 - `init()`'s all libraries
 - `_` then calls `main()`
 - ...and `exit()`
- Sometimes loaders are called “runtime linkers”

Interpreter

`readelf --headers /bin/ls`

Lazy binding

- Binding a function call to a library can be expensive
 - Have to go through code and replace the symbol with its address
- Delay until the call actually takes place
 - Calls stub PLT function
 - Invokes dynamic linker to load the function into memory and obtain real address
 - Rewrites address that the sub code references
 - Only happens once
- Procedure Lookup Table (PLT)

Questions?