



CS 25200: Systems Programming

Lecture 6: Memory Management and malloc()

Prof. Turkstra



Lecture 06

- `brk()` and `sbrk()`
- Free lists
- Fragmentation
- `malloc()` internals
- Segregated free lists
- Coalescing
- Fenceposts

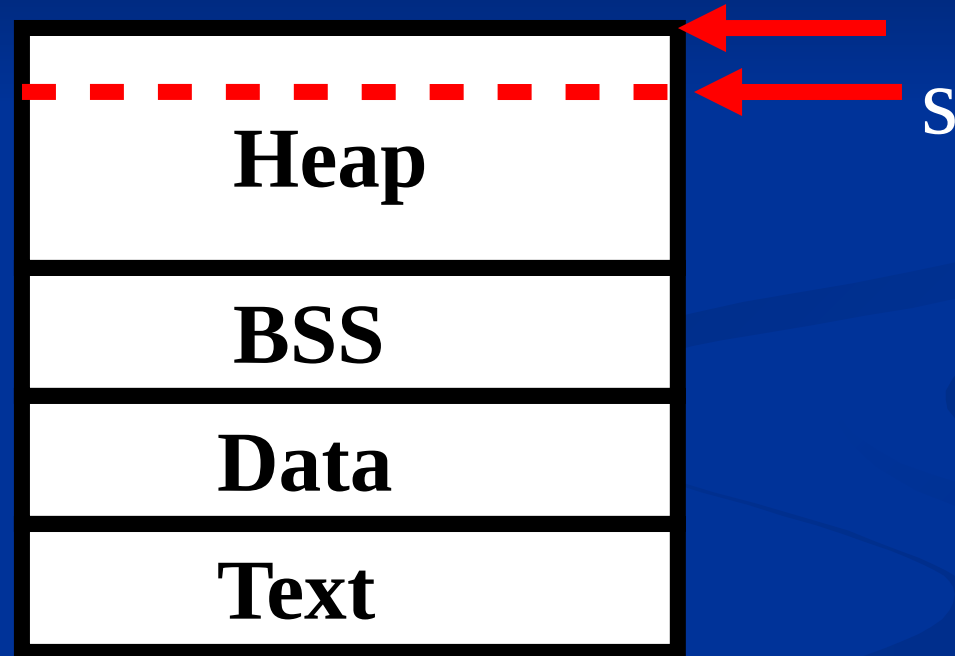
The Heap

- From the kernel's standpoint, the heap is a single contiguous region of memory that grows linearly
- To request more memory, a userland process invokes the `brk()` system call
 - libc provides two wrapped versions of `brk()`: `sbrk()` and `brk()`

brk() vs sbrk()

- `int brk(void *addr);`
Set the end of the data segment to the value specified by *addr* when *reasonable*
 - Cannot exceed maximum data size
- `void *sbrk(intptr_t increment);`
Increment the program's data space by *increment* bytes.
 - Returns the previous program break or `(void *) -1`
 - `sbrk(0)` returns the current location
 - Argument can be negative

`s = sbrk(n);`



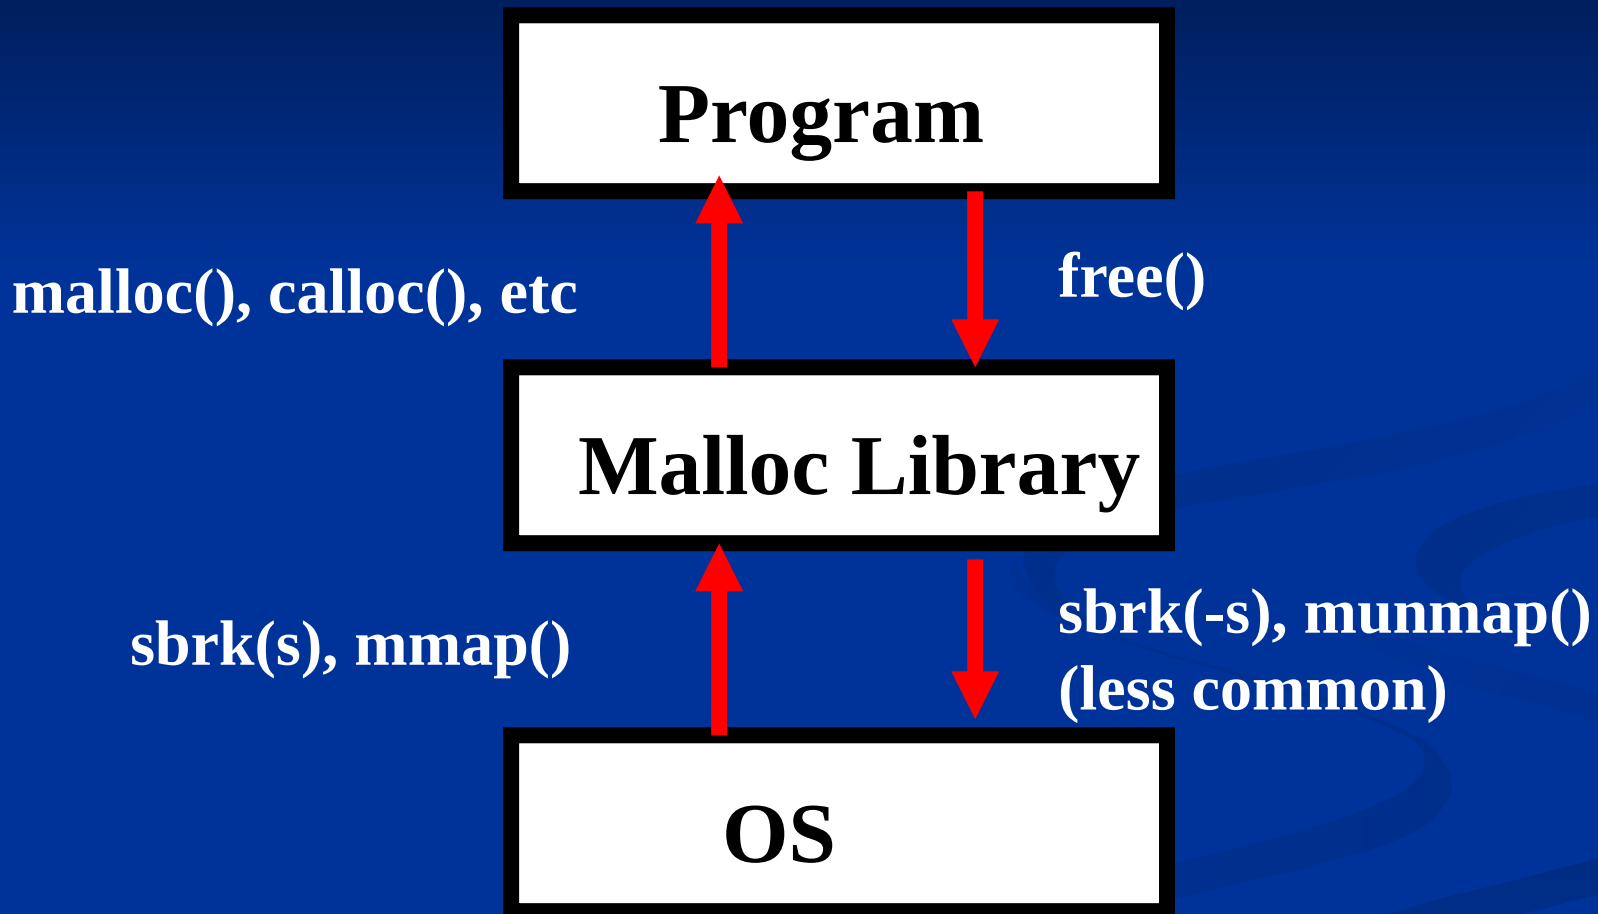
malloc

- The portable and “comfortable” way to allocate memory in C is by using the provided **memory allocation package**
 - malloc(), calloc(), realloc(), and free()
- The program break is managed for you, internally
- Programmer explicitly invokes the above functions

How?

- Memory is requested from the OS in large “chunks” (e.g., 64KiB)
- These chunks are then managed internally
 - Added to a **free list**
 - Subsequent requests are satisfied from the free list when possible
- Decreases number of times the OS must be invoked (via system call)

Memory management



Implementation

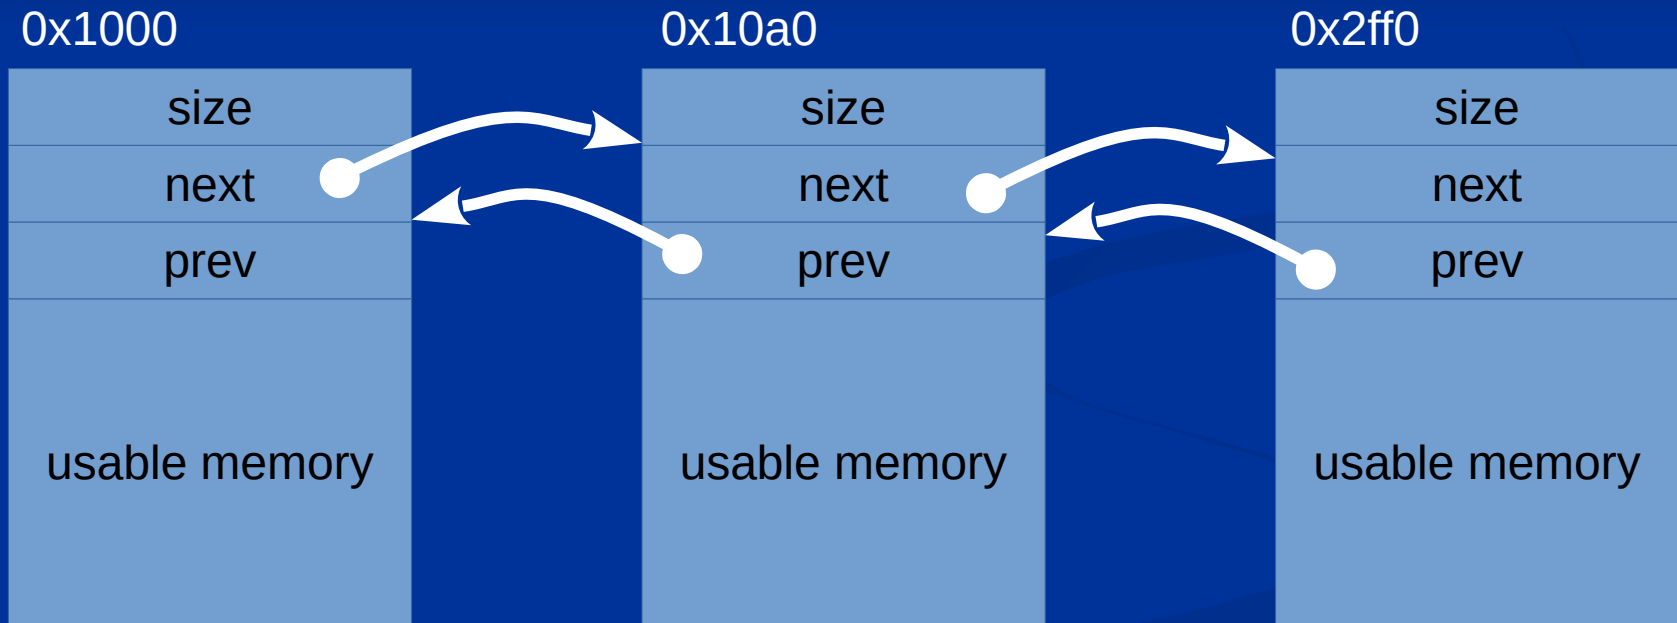
- There are many different ways to implement allocators
- Some useful data structures...
 - Single free list
 - Segregated free lists
 - Cartesian trees
 - Boundary tags

Single free list

- ...or sequential fit
- Structure that is sequentially searched to find the needed size
 - First-fit
 - Best-fit
 - Next-fit
 - Worst-fit
- Can be a singly-linked list, doubly-linked list, tree, etc

Where is the list stored?

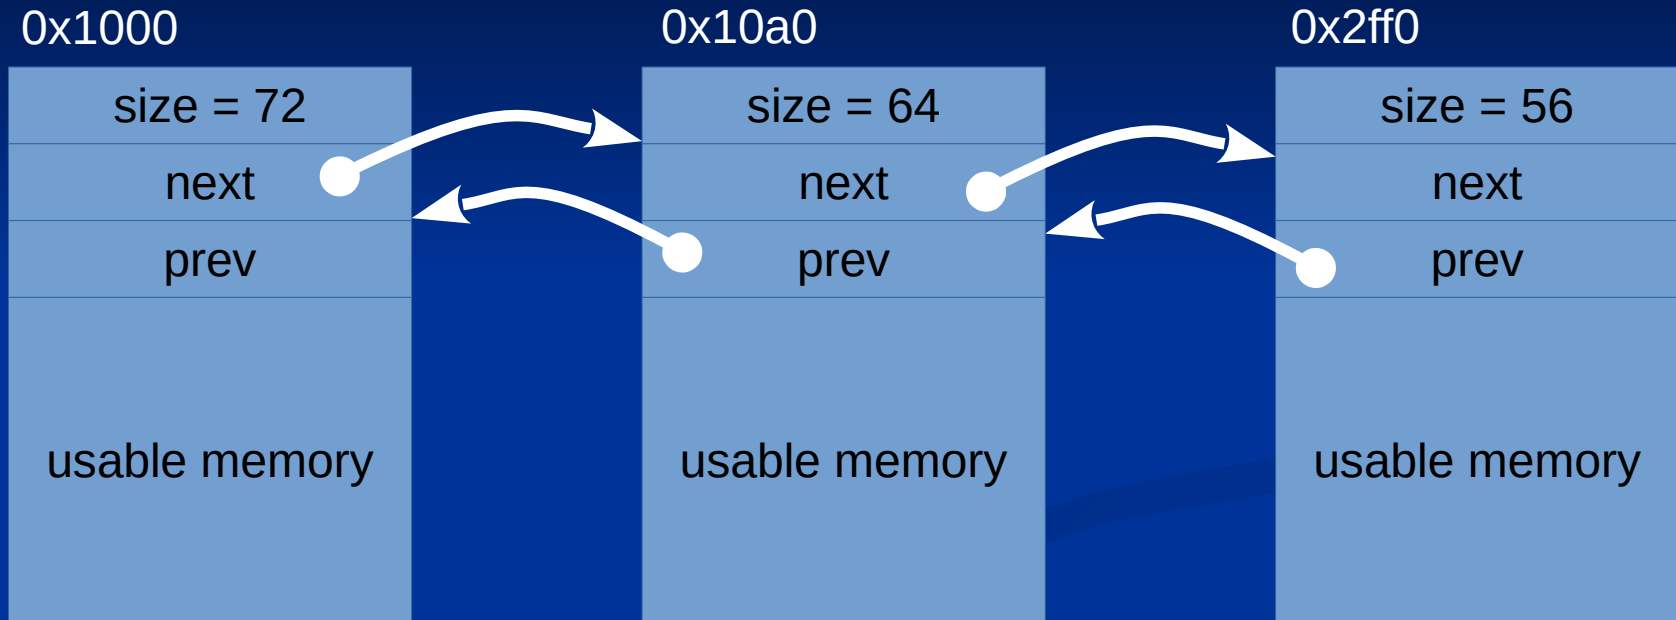
- In the free blocks!
- Each block has a header



Fragmentation

- Memory that is too small to be usefully allocated
 - External: visible to allocator
 - Internal: visible to requester
- Want to minimize fragmentation

External fragmentation



```
p = malloc(100);
```

- 104 + 16 bytes (header) = 120 bytes total

```
p = ??
```

External fragmentation

- Can be calculated:

$$\text{Ext_Frag} = 100 * (1 - \text{size}(\text{largest_free_block}) / \text{sum}(\text{free_mem}))$$

- Previous example:

$$100 * (1 - 72 / (72 + 64 + 56)) = 62.5\%$$

- Only one block? 0%

Comparison

- Fragmentation depends on the algorithm and the **workload**
- Best fit tends to leave some very large holes and some very small holes
 - Can't use small holes easily
 - Computationally more expensive
- First fit tends to leave “average” sized holes
 - Also faster than best fit
- Next fit often used in practice
 - Prevents accumulation of small chunks at the beginning

Mechanics


- `malloc()`: search the free list for the appropriate size
 - Found? Split if necessary (add the remainder to free list) and return the block
 - Not found? Request more memory (chunk) from the OS, add to free list, repeat
 - Usually done in “large” increments (e.g., 4KiB)

free()

- Coalesce if possible
 - If not, insert into list
- How does free() know the size of the memory chunk passed to it?

Header

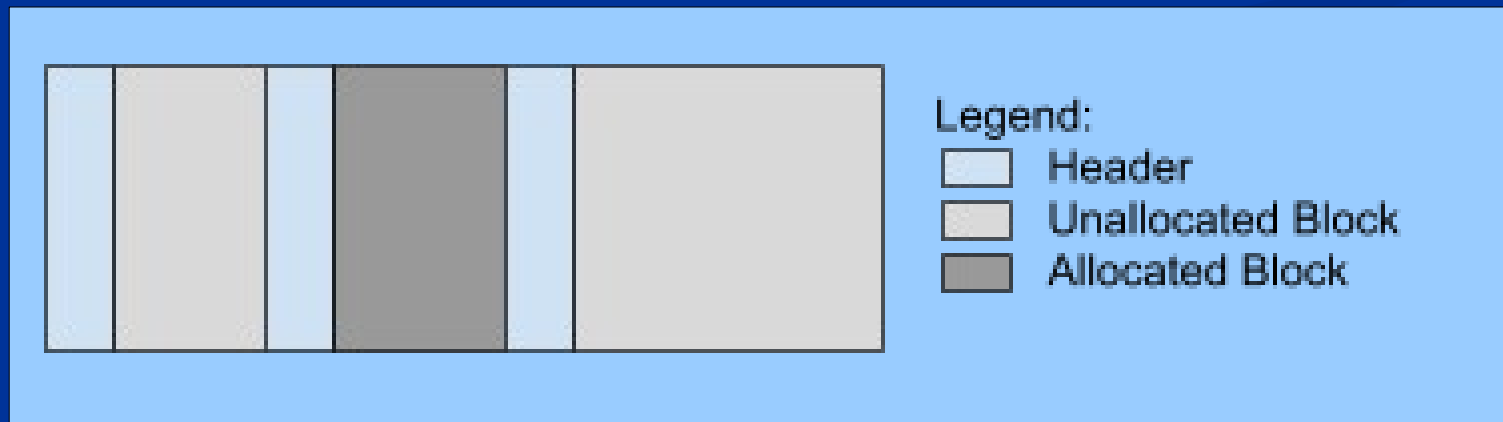
- Remember the free list? malloc()'d chunks also have a header
 - But no list



- malloc() returns a pointer that points **after** the header

Headers

- Easy to get the next block header
(header *) ((char *) block_addr +
ALLOC_HEADER_SIZE + size)
- What about the previous block?



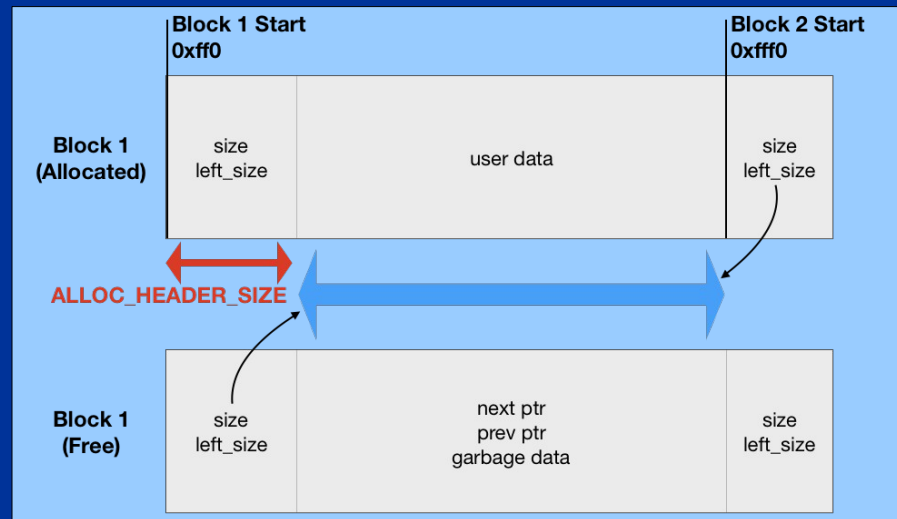
Boundary tag

- Track the size of the “left” block
 - ...in the current block’s header
- Previous block header is now:
$$(\text{header} *) ((\text{char} *) \text{block_addr} - \text{left_size} - \text{ALLOC_HEADER_SIZE})$$


Constant time coalescing

- Donald Knuth

- Note that both size and left_size include the next and previous pointers



- This is the maximum possible usable memory when allocated



How to tell allocation status?

Naive Solution	Our Solution (Allocated)	Our Solution (Unallocated)
size	size allocated	size allocated
left_size	Left_size	left_size
allocated		next
next	User's data	prev
prev		Undefined data
User's data Undefined data		

Bitwise operations!

- How do you set a bit?
`size |= 0x1`
- How do you clear them?
`size &= ~0b111`

Header

```
typedef struct header {  
    size_t size;  
    size_t left_size;  
    union {  
        struct {  
            struct header *next;  
            struct header *prev;  
        };  
        char *data;  
    }  
} header;
```

Alignment

- Many RISC architectures simply cannot handle an unaligned access
 - Sparc: SIGBUS
- x86 can, but it is **slow**
- Our malloc() should always return a MIN_ALLOCATION-aligned address
 - 8 bytes for now, could change!

Minimum allocated size

- Suppose the user requests 1 byte:
`malloc(sizeof(char));`
 - Must be a multiple of
`MIN_ALLOCATION`, so round to 8 bytes
 - Good enough?

Internal fragmentation


- Waste of memory due to allocator returning a larger block than requested

$\text{Int_Frag} = 100 * (1 - \text{size}(\text{request}) / \text{sum}(\text{mem_allocated}))$

- E.g., $\text{malloc}(1) \rightarrow 8 + \text{roundup8}(1)$
= 40 bytes for 1 byte of memory

Segregated free lists


- Multiple free lists, one for each size



- Objects allocated from the free list of nearest size
- Empty or no size large enough? Get more memory from the OS and populate the appropriate free list

Coalescing

- Some implementations do not coalesce
 - Ours will
- Without a footer or boundary tag, requires traversal to find out if neighboring blocks are free



Segregated lists


- With segregated lists,
 - Allocation is often $O(1)$
 - `free()` would be $O(1)$ without coalescing
- Segregated free lists are fast
 - But use more memory
 - Even more memory if they don't coalesce

Project 2

- You will implement a malloc library that can be used as a substitute for libc's malloc
- Gain an understanding of malloc internals
- Also better understand memory errors – premature free, double free, wild free, etc

Fenceposts

- Memory inside malloc is obtained a “chunk” at a time
- Chunks can also be coalesced
- How do we know when not to?




Obtaining chunks

- Round to nearest `ARENA_SIZE`
 - Don't forget to include size of fenceposts and header in request
- Call `sbrk()`
 - Determine if the new request is contiguous with previous `brk`
 - Yes? Eliminate fencepost
 - No? Don't
- Initialize fenceposts (`set_fenceposts()`)
 - May have to fix `left_size`
- Initialize header
- Add to free list
 - Possibly coalesce



All together



Fencepost
Header
Footer
Free
Allocated

Questions?