



CS 25200: Systems Programming

Lecture 8: Memory Allocation Wrap-up Program Generation and Loading

Prof. Turkstra



Lecture 08

- Memory allocation wrap-up
- Executable formats
- Building a program
- Libraries
- Loader

Memory smashing

- Program writes to memory outside of the intended data structure or allocated range
- Corrupts the following object(s)
 - In the heap, may overwrite boundary tags or even fencepost(s)
- Sometimes nothing happens
 - Why?

Example

```
char *s = malloc(8);  
strcpy(s, "Hello, World!");
```

Debugging memory allocation errors

- Can be very difficult
 - Corruption of malloc's internals may not become apparent until much later
 - Memory leaks may not be immediately obvious
- Trying to find premature/double/wild frees?
 - Can try commenting out all free() calls
 - If the problem goes away, uncomment them one-by-one

Debugging

- ...or use a tool
- Free
 - Valgrind
 - Memcheck
- Not free
 - IBM Rational Purify
 - Bounds Checker
 - Insure++

Tradeoffs

- Explicit memory allocation is...
 - Fast
 - Efficient with regard to memory usage
- It also...
 - Can be error prone
 - Requires more expertise
 - May take longer to develop and debug
 - Can be insecure if not done right

- Some prefer to use languages like Java or C#
- C/C++ is still widely used
 - Especially for lower-level, performance critical applications

Debugging allocators

- Can always augment your allocator to make it more robust
- Check if a block was returned by malloc() using a magic number in the header
 - Prevents wild free()s
- Use different magic numbers to indicate free/allocated
 - Instead of a single bit
 - E.g. free: 0xF7EEF7EE; allocated: 0xA10CA7ED
 - Always check the magic number

- Create an integrity checking function
 - Iterate over all free and allocated blocks
 - Check for two consecutive free blocks
 - Ensure there are no unexpected gaps
- Store size of gaps in the fence posts
- Verify sizes and flags in header match
- Print free and allocated blocks

Gaps

- Often gaps between segments
- Attempting to access an address in an **unmapped** region causes the OS to send the process a signal: **SIGSEGV**
 - Segmentation violation
 - By default program is immediately terminated and dumps **core**
- This also happens if you access a mapped but **protected** region
 - E.g., try executing in the data segment

Core dump file

- `$ man 5 core`
- File containing an image of the process' memory at time of termination
- Can be used with gdb (or other debuggers)
- May have to enable it (e.g., on `data.cs.purdue.edu`)
 - `$ ulimit -c unlimited`



Program

- File in a particular format containing necessary information to load an application into memory and execute it
 - Often time part of this is split off into the “loader” and libraries
- Programs include:
 - Machine instructions
 - Initialized data
 - List of library dependencies
 - List of memory sections
 - List of values determined at load time

Executable file formats

- Fully “compiled” programs can come in many different formats...
 - ELF – Executable Linux File
 - Used by most recent UNIX systems (e.g., Solaris, Linux)
 - PE – Portable Executable
 - Used by Windoze
 - Mach-O – Mach object [file format]
 - Used by macOS/iOS
 - COFF – Common Object File Format
 - Also windoze, some embedded systems - historically System V Unix
 - a.out – Used in BSD (Berkeley Standard Distribution)
 - Restrictive, rarely seen anymore

ELF

- File header
 - Magic number
 - Version
 - Target ABI
 - ISA
 - Entry point
 - Pointers to
 - Program header
 - Section header
 - etc

Program header

- How to create the process image
 - Segments
 - Types
 - Flags
 - File offset
 - Virtual address
 - Size in file
 - Size in memory

Section header

- Type (data, string, notes, etc)
- Flags (writable, executable, etc)
- Virtual address
- Offset in file image
- Size
- Alignment

- `readelf --headers /bin/ls`
- `objdump -p, -h, -t`

Purdue "All-American" Marching Band

- Will be performing tomorrow starting at 10:30 AM, Slayter Center



Building a program

- Start with source code
 - abs.c
- Preprocessor
- Compiler
- Assembler
- Linker




Figure 4.6 The steps used to translate a source program to the binary object code representation used by a processor.

Preprocessor

- When a .c file is compiled, it is first scanned and modified by the preprocessor before being handed to the real compiler
- Finds lines beginning with #, hides them from the compiler, or takes some action
- #include, #define
- #ifdef, #else, #endif

Why macros?

- Run time efficiency
 - No function call overhead
- Passed arguments can be any type
 - `#define MAX(x,y) ((x) > (y) ? (x) : (y))`
 - Works with ints, floats, doubles, even chars

■ Can do math

- `#if (FLAG % 4 == 0) || (FLAG == 13)`

■ Macros

- `#define INC(x) x+1`
- No semi-colon
- Have to be careful
 - `#define ABS(x) x < 0 ? -x : x`
 - `ABS(B+C)`

- Parentheses around substitution variables

```
#define ABS(x) ( (x) < 0 ? -(x) : (x) )
```


Lots of other tricks

```
printf("The date is %s\n", __DATE__);
```

- Most preprocessor features are used for large/advanced software development practices

- `gcc -E file.c > output.i`

Compiler?



* <http://www.cs.montana.edu/~david.watson5/>

■ gcc -S

Assembler

- Discussed in CS 250 – Computer Architecture
- `gcc -c`
`nm -v`
- Really uses `as`

Libraries

- Libraries are just collections of object files
 - Internal symbols are indexed for fast lookup by the linker
- Searched for symbols that aren't defined in the program
 - Symbol found, pull it into executable (static)
 - Otherwise include a pointer to the file, loaded by loader

Statically linked

- Faster, to a degree
- Portable
- Larger binaries
- Fixed version, no updates
- File extension .a

Dynamically linked

- More complexity
- Easy to upgrade libraries
 - Vulnerabilities
- Have to manage versions
- Loader re-links every time program is executed

```
readelf --dynamic /bin/ls  
ldd /bin/ls
```

End result

- `gcc -o abs abs.c`
`nm -v ./abs`

Loader

- Essential step in starting a program
- Historically allocated space for all sections of the executable (text, data, bss, etc)
- Now simply establishes mappings
 - Page faults actually populate the memory
 - For executable as well as (shared) libraries

- Also resolves any values in the executable to point to the functions/variables in the shared libraries
- Jumps to `_start`
 - `init()`'s all libraries
 - `_` then calls `main()`
 - ...and `exit()`
- Sometimes loaders are called “runtime linkers”

Interpreter

`readelf --headers /bin/ls`

Lazy binding

- Binding a function call to a library can be expensive
 - Have to go through code and replace the symbol with its address
- Delay until the call actually takes place
 - Calls stub PLT function
 - Invokes dynamic linker to load the function into memory and obtain real address
 - Rewrites address that the sub code references
 - Only happens once
- Procedure Lookup Table (PLT)

Questions?