



## **CS 25200: Systems Programming**

### **Lecture 14: Shell Executor, Processes, and Signals**

Prof. Turkstra



# Voting

- Out of state IDs are not accepted
- Federal IDs are fine (passport, military, etc)
- IDs must have an expiration date
  - First year this is enforced
- You can get a replacement Purdue ID for free October 21 - 25
  - Otherwise it is \$10



# Final Exam

- Tuesday, December 10 1pm – 3pm,  
LILY 1105



# Midterm Exam

- Thursday, October 17 8pm - 10pm  
WALC 1055
- Make up lecture on Monday 11/25?
  - Week of Thanksgiving break

# Lecture 13

- Shell executor
- `exit()` vs `_exit()`
- Processes
- Signals

# Another look at ls\_grep.c


- Included in Lab 3's example directory

# FD\_CLOEXEC

- Causes file descriptor to be closed any time a call is made to a function in the exec() family
- Flag is stored in the file descriptor table, not the file object
  - Why?

# Executor

- Parent process sets up all piping and redirection before forking children
- Children inherit redirections
- Parent must save original fds and restore them
- stderr should be the same for all processes





```

execute() {
    // save in/out
    int default_in = dup(0);
    int default_out = dup(1);

    // set the initial input
    int input_fd;
    if (input_file) {
        input_fd = open(input_file,.....);
    }
    else {
        // Use default input
        input_fd = dup(default_in);
    }

    int ret;
    int output_fd;
    for (i = 0; i < num_single_commands; i++) {
        // redirect input
        dup2(input_fd, 0);
        close(input_fd);

        // setup output
        if (i == num_single_commands - 1){
            // Last single command
            if (output_file) {
                output_fd = open(output_file,.....);
            }
        }
    }
}

```



```

        else {
            // Use default output
            output_fd = dup(default_out);
        }
    }
    else {
        // Not last single command
        // Create pipe
        int pipe_fds[2];
        pipe(pipe_fds);
        output_fd = pipe_fds[1];
        input_fd = pipe_fds[0];
    }

    // Redirect output
    dup2(output_fd, 1);
    close(output_fd);

    // Create child process
    ret = fork();
    if (ret == 0) {
        close(default_in);
        close(default_out);
        execvp(scmd[i].args[0], scmd[i].args);
        perror("execvp");
        exit(1);
    }
}

```

```
// restore in/out defaults
dup2(default_in,0);
dup2(default_out,1);
close(default_in);
close(default_out);

if (!background) {
    // Wait for last command
    waitpid(ret, NULL,0 );
}

} /* execute() */
```

# exit() and \_exit()

- **void exit(int status)** – cause normal process termination
  - All atexit() and on\_exit() registered functions are executed
    - In reverse order
  - stdio streams are flushed and closed
  - tmpfile() files are removed
- **void \_exit(int status)** – terminate the process immediately
  - fds still closed

# Shell strategy notes

- Remember, `input_fd` tracks the input for the **next** command
  - Will either be the input file fd (for the first command) or `pipe_fds[0]`
- Example only handles pipes and input/output redirection
- Have to handle `stderr` redirection
  - Applies to all processes
- Also have to handle the “append” cases

# Shell

Final Command Table

ls	-al	aab	aaa
grep	me		
In:dflt	Out:file1	Err:dflt	

Lexer

shell.l

Parser

shell.y

wildcards  
env vars

executor

ls -al a\* | grep me > file1

<ls> <-al>  
<a\*> <PIPE>  
<grep> <me>  
<GREAT>  
<file1>

Command Table

ls	-al	a*
grep	me	
In:dflt	Out:file1	Err:dflt

# Program vs. process

- A **program** is an executable file that contains a set of instructions
  - Usually stored on disk or other secondary storage
- A **process** is a program in execution
  - It resides, at least partially, in memory

# Processes

- Programs may have multiple processes or instances running
  - E.g., multiple instances of Bash
- All processes have a parent
  - Except init, pid 1
- Remember ps?



# Processes

- top, ps (-e, -ax, -f, -u)

# Process properties

- PID: Process ID, index into process table
- Command/program name
- Arguments
- Environment variables
- Current working directory
- User ID
- stdin / stdout / stderr

# Process ID

- Uniquely identifies running process
- Initial process (init, systemd) has PID 1
- PIDs assigned in ascending order
- Wrap around when limit is reached
- System call to get pid:  
`pid_t getpid();`

# Command and arguments

- Every process has a command name and 0 or more arguments
- Arguments are passed to main  

```
int main(int argc, char **argv);
```

  - `argc`: number of args
  - `argv`: arguments (`argv[0]` is the command name)

# printargs.c

```
int main(int argc, char **argv) {  
    for (int i = 0; i < argc; i++) {  
        printf("argv[%d]=\"%s\"\\n", i, argv[i]);  
    }  
}
```

# Environment variables

- Array of strings, A=B, inherited from the parent process
- E.g.
  - PATH=/bin:/usr/bin - directories to search for commands
  - USER=<login> - username
  - HOME=/homes/turkstra
- Can modify .login or .bashrc
  - Aliases, etc too

# Manipulating

- `export A=B`
  - All children will also see the change
- `A=B`
  - Only the current process will get it
- For example,  
`export PATH=$PATH:~/bin`
- Can run `env` or `export` to view current environment

# In C

```
extern char **environ;
```

```
int main(int argc, char **argv) {  
    int i = 0;  
    while (environ[i] != NULL) {  
        printf("%s\n", environ[i]);  
        i++;  
    }  
}
```



# Current directory

- Sometimes called working directory, current working director, or present working directory
  - `$ pwd`
- Every process has a current directory
  - Really just an inode

- Used to resolve relative paths
- Relative paths do not begin with /
  - /root/hello.c – absolute
  - hello.c – relative
  - ../src/hello.c – relative
  - src/hello.c – relative
  - ~/src/hello.c – relative
- Change current directory: `$ cd dir`
- System call:  
`int chdir(const char *path);`  
`int fchdir(int fd);`

# User identifier

- Processes have an effective user ID (euid)
  - Files created are owned by it
  - Most access checks use it
- ...and a real uid (ruid)
  - Inherited from parent
  - Impacts signal sending/receiving

- Processes running as root can change their UID using:

```
int setuid(uid_t uid);
```

- It's permanent.
- E.g.: OpenSSH runs as root, but setuid()s to the connecting user after fork.

- Or...

```
int seteuid(uid_t uid);
```

- Allows setuid-root processes to temporarily drop root privileges

- Remember **sudo** and **su**?



# Signals

- One form of inter-process communication (IPC)
- Asynchronous mechanism for the OS to communicate with a running process
- Processes can register signal handlers to perform certain actions for certain signals
- Signals are similar to interrupts

# Some signals

- SIGHUP: Hangup
- SIGINT: Terminal interrupt
- SIGBUS: BUS error
- SIGKILL: Kill (cannot be ignored)
- SIGSEGV: Segmentation violation
- SIGTERM: Termination
- SIGCHLD: Child process has stopped, exited, or changed
- SIGUSR1, SIGUSR2, etc

# Handling signals

- `sighandler_t signal(int signum,  
                          sighandler_t handler)`
- `int sigaction(int signum,  
              const struct sigaction *act,  
              struct sigaction *oldact);`

```
struct sigaction {  
    void      (*sa_handler)(int);  
    void      (*sa_sigaction)(int, siginfo_t *, void *);  
    sigset_t   sa_mask;  
    int        sa_flags;  
    void      (*sa_restorer)(void);  
};
```

# Flags

- **SA\_RESTART**: Resume the function after a signal is handled properly
- Instead of returning **EINTR**
- **SA\_NOCLDSTOP**: Only deliver **SIGCHLD** on termination, not stopping
- **SA\_ONSTACK**: Use the signal stack
  - Must set it up first



# Signals and lex

- Lex's scanner uses `getc()` to read from fd 0
- `getc()` is built on top of the `read()` system call
- Many blocking system calls will return if a signal is received
  - And set `errno` to `EINTR`
- What happens when we get `SIGINT` (or `SIGCHLD`)?
  - `getc()` returns -1!
- How do we stop it?

# Keeping lex alive

- ...use SA\_RESTART

```
struct sigaction signal_action;
```

```
signal_action.sa_handler = sig_int_handler;
```

```
sigemptyset(&signal_action.sa_mask);
```

```
signal_action.sa_flags = SA_RESTART;
```

```
int error = sigaction(SIGINT,  
                      &signal_action, NULL);
```

```
if (error) {  
    perror("sigaction");  
    exit(-1);  
}
```

# Questions?