



CS 25200: Systems Programming

Lecture 15: Signals, Built-ins, and Subshells

Prof. Turkstra



Evaluations

- Mid-semester course evaluations available
- Some notes on third party “evaluations”

2b - University questions about your instructor		Turkstra, Jeffrey A								
		Responses					Individual			
		E	G	F	P	VP	N	Grp Med	Mode	Std Dev
Q6	Overall, I would rate this instructor as:	17	36	13	7	4	77	3.9	4	1.07
Responses: [E] Excellent=5 [G] Good=4 [F] Fair=3 [P] Poor=2 [VP] Very Poor=1										

3c - PICES instructor based questions		Turkstra, Jeffrey A								
		Responses					Individual			
		SA	A	U	D	SD	N	Grp Med	Mode	Std Dev
Q19	My instructor takes my views and comments seriously.	21	31	18	3	1	74	4.0	4	.90
Responses: [SA] Strongly Agree=5 [A] Agree=4 [U] Undecided=3 [D] Disagree=2 [SD] Strongly Disagree=1										

CS standard questions about instructor		Turkstra, Jeffrey A								
		Responses					Individual			
		SA	A	U	D	SD	N	Grp Med	Mode	Std Dev
Q20	My instructor seemed concerned that students learn.	21	35	13	4	1	74	4.0	4	.89
Q21	My instructor showed a clear understanding of the subject.	45	22	5	0	2	74	4.7	5	.84
Q22	My instructor was well-prepared and organized in class.	30	32	11	1	0	74	4.3	4	.75
Q23	My instructor helped me understand the material.	13	32	20	7	2	74	3.8	4	.97
Q24	My instructor gave exams that accurately assessed what I have learned in this course.	14	28	13	13	6	74	3.7	4	1.21
Q25	My instructor graded fairly.	15	33	17	8	0	73	3.8	4	.90
Q26	My instructor was reasonably available to help students outside of class.	20	35	14	3	0	72	4.0	4	.80
Q27	My instructor encouraged in-class participation.	25	36	10	3	0	74	4.2	4	.79
Q28	My instructor shows respect for me and other students in this class.	23	35	12	3	1	74	4.1	4	.87
Q29	My instructor is open to student's questions.	36	33	4	1	0	74	4.5	5	.66
Q30	My instructor encourages an atmosphere where ideas can be exchanged freely and easily.	30	34	7	2	1	74	4.3	4	.83
Responses: [SA] Strongly Agree=5 [A] Agree=4 [U] Undecided=3 [D] Disagree=2 [SD] Strongly Disagree=1										

Grades

- Quiz grades updated
 - Quiz 9 grade
- Lab attendance up soon
- Remember, regrade deadline is **one week**

Project 1 Stats

- Average: 105.5/125 (84.4%)
- Standard Deviation: 22.7/125 (18.2%)
- Median: 114.5/125 (91.6%)

Project 2 Stats

- Average: 78.7
- Standard Deviation: 28.9
- Median: 91.6

* excluding students that did not submit

Lecture 15

- Signals
- Built-ins
- Subshells

Signals

- One form of inter-process communication (IPC)
- Asynchronous mechanism for the OS to communicate with a running process
- Processes can register signal handlers to perform certain actions for certain signals
- Signals are similar to interrupts

Some signals

- SIGHUP: Hangup
- SIGINT: Terminal interrupt
- SIGBUS: BUS error
- SIGKILL: Kill (cannot be ignored)
- SIGSEGV: Segmentation violation
- SIGTERM: Termination
- SIGCHLD: Child process has stopped, exited, or changed
- SIGUSR1, SIGUSR2, etc

Handling signals

- `sighandler_t signal(int signum,
 sighandler_t handler)`
- `int sigaction(int signum,
 const struct sigaction *act,
 struct sigaction *oldact);`

```
struct sigaction {  
    void      (*sa_handler)(int);  
    void      (*sa_sigaction)(int, siginfo_t *, void *);  
    sigset_t   sa_mask;  
    int       sa_flags;  
    void      (*sa_restorer)(void);  
};
```

Flags

- **SA_RESTART**: Resume the function after a signal is handled properly
- Instead of returning **EINTR**
- **SA_NOCLDSTOP**: Only deliver **SIGCHLD** on termination, not stopping
- **SA_ONSTACK**: Use the signal stack
 - Must set it up first

Signals and lex

- Lex's scanner uses `getc()` to read from fd 0
- `getc()` is built on top of the `read()` system call
- Many blocking system calls will return if a signal is received
 - And set `errno` to `EINTR`
- What happens when we get `SIGINT` (or `SIGCHLD`)?
 - `getc()` returns -1!
- How do we stop it?

Keeping lex alive

- ...use SA_RESTART

```
struct sigaction signal_action;
```

```
signal_action.sa_handler = sig_int_handler;
```

```
sigemptyset(&signal_action.sa_mask);
```

```
signal_action.sa_flags = SA_RESTART;
```

```
int error = sigaction(SIGINT,  
                      &signal_action, NULL);
```

```
if (error) {  
    perror("sigaction");  
    exit(-1);  
}
```

Shell

Final Command Table

ls	-al	aab	aaa
grep	me		
In:dflt	Out:file1	Err:dflt	

Lexer

shell.l

Parser

shell.y

wildcards
env vars

executor

ls -al a* | grep me > file1

<ls> <-al>
<a*> <PIPE>
<grep> <me>
<GREAT>
<file1>

Command Table

ls	-al	a*
grep	me	
In:dflt	Out:file1	Err:dflt

Built-ins

- Some “commands” are not discrete programs
- Usually simple ways to alter the shell’s state
 - E.g. `cd`, `setenv`, `unsetenv`, `printenv`, etc

Built-ins

- Should we fork?
- cd: See the chdir() system call
- setenv: Set an environment variable
- unsetenv: Unset it
- printenv: Display environment variables
- source: Execute each line as though it were in the input stream

printenv

- Linux implements printenv as a program :-)
- But...

```
ret = fork();
if (ret == 0) {
    if (!strcmp(argument[0], "printenv")) {
        char **p = environ;
        while (*p != NULL) {
            printf("%s", *p++);
        }
        exit(0);
    }
    ...
    execvp(...);
}
```

Velociraptors

- "It's a UNIX system! ... I know this!"




Subshell

- Subshells allow commands to execute in a shell environment without impacting the parent shell
 - Copy of output is returned to the parent shell
 - Parent shell fork()s and exec()s itself
 - Requires bidirectional communication

Two pipes

```
int pin[2];  
int pout[2];
```

```
pipe(pin);  
pipe(pout);
```



- Redirect **pin[0]** to child's input
- Redirect child's output to **pout[1]**

- Parent writes the command to `pin[1]`
 - And some other things – what?
- Can then read output from `pout[0]`
 - This is then fed back into the lexer
 - In reverse order...
 - See `myunputc()`

Shell

Final Command Table

ls	-al	aab	aaa
grep	me		
In:dflt	Out:file1	Err:dflt	

Lexer

shell.l

Parser

shell.y

wildcards
env vars

executor

ls -al a* | grep me > file1

<ls> <-al>
<a*> <PIPE>
<grep> <me>
<GREAT>
<file1>

Command Table

ls	-al	a*
grep	me	
In:dflt	Out:file1	Err:dflt

Wildcards

- Allow us to perform actions on more than one file at a time
- Also called file globbing
- Do not confuse with regular expressions

Standard wildcards

- `$ man 7 glob`
 - `?`: any single character
 - `DSCN00??.JPG`
 - `*`: any number of characters, including zero
 - `a*a *.c t*.log`
 - `[]`: a range
 - `m[a-d]m m[0-7]n`
 - `{ }`: “or” relationship – no spaces allowed
 - `cp {*.doc,*.pdf} ~`
 - `[!]`: anything not in the range
 - `\`: escape character

Our shell

- Only need to worry about * and ?
- Implement the simple case first
 - Add a function to shell.y to do the expansion

shell.y

■ Before...

```
argument: WORD {  
    insert_single_command(current_command, $1);  
} ;
```

■ After...

```
argument: WORD {  
    expand_wildcards($1);  
} ;
```

POSIX Regular expressions

- `int regcomp(regex_t *preg, const char *regex, int cflags)`
 - Compiles the regular expression into a form that `regexexec()` can use
 - Completes the passed `regex_t`
 - Finite state automata
- `int regexexec(const regex_t *preg, const char *string, size_t nmatch, regmatch_t pmatch[], int eflags)`
 - Do the actual comparison

```
typedef struct {  
    regoff_t rm_so;  
    regoff_t rm_eo;  
} regmatch_t;
```

void regfree(regex_t *preg)

- regex_t's contain dynamically allocated memory
- You have to free it
 - ...or you'll have a memory leak
- Even have to free before reusing!
- Does **not** free the regex_t itself

Questions?