



CS 25200: Systems Programming

Lecture 19: Processes and Scheduling

Prof. Turkstra



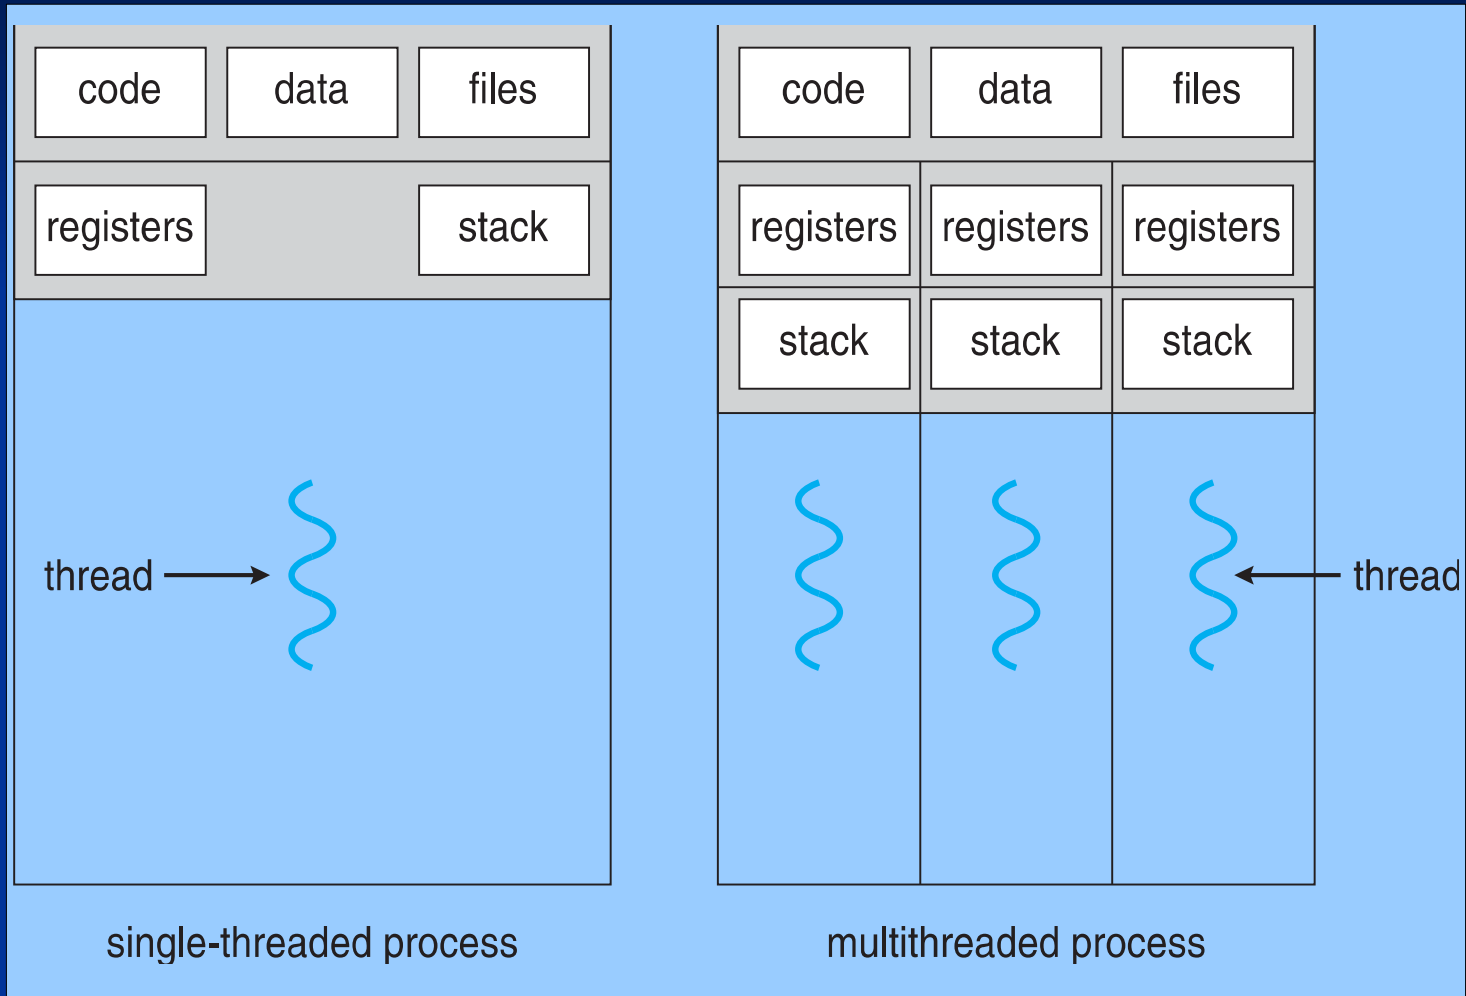
Lecture 19

- Threads and context switches
- Non-preemptive vs preemptive scheduling
- FCFS scheduling
- Round robin scheduling
- Multilevel feedback queue scheduling

Threads

- Process includes...
 - Address space (code, data, etc)
 - Resource container (OS resource, accounting)
 - A “thread” of control – PC, regs, stack
- Threads
 - Share some code and data (address space)
 - Same files, I/O channels, resource containers
 - Do **not** share thread of control

Threads



Threads

- Can have several threads in a single address space
- Threads are units of scheduling
- Processes are containers in which threads execute

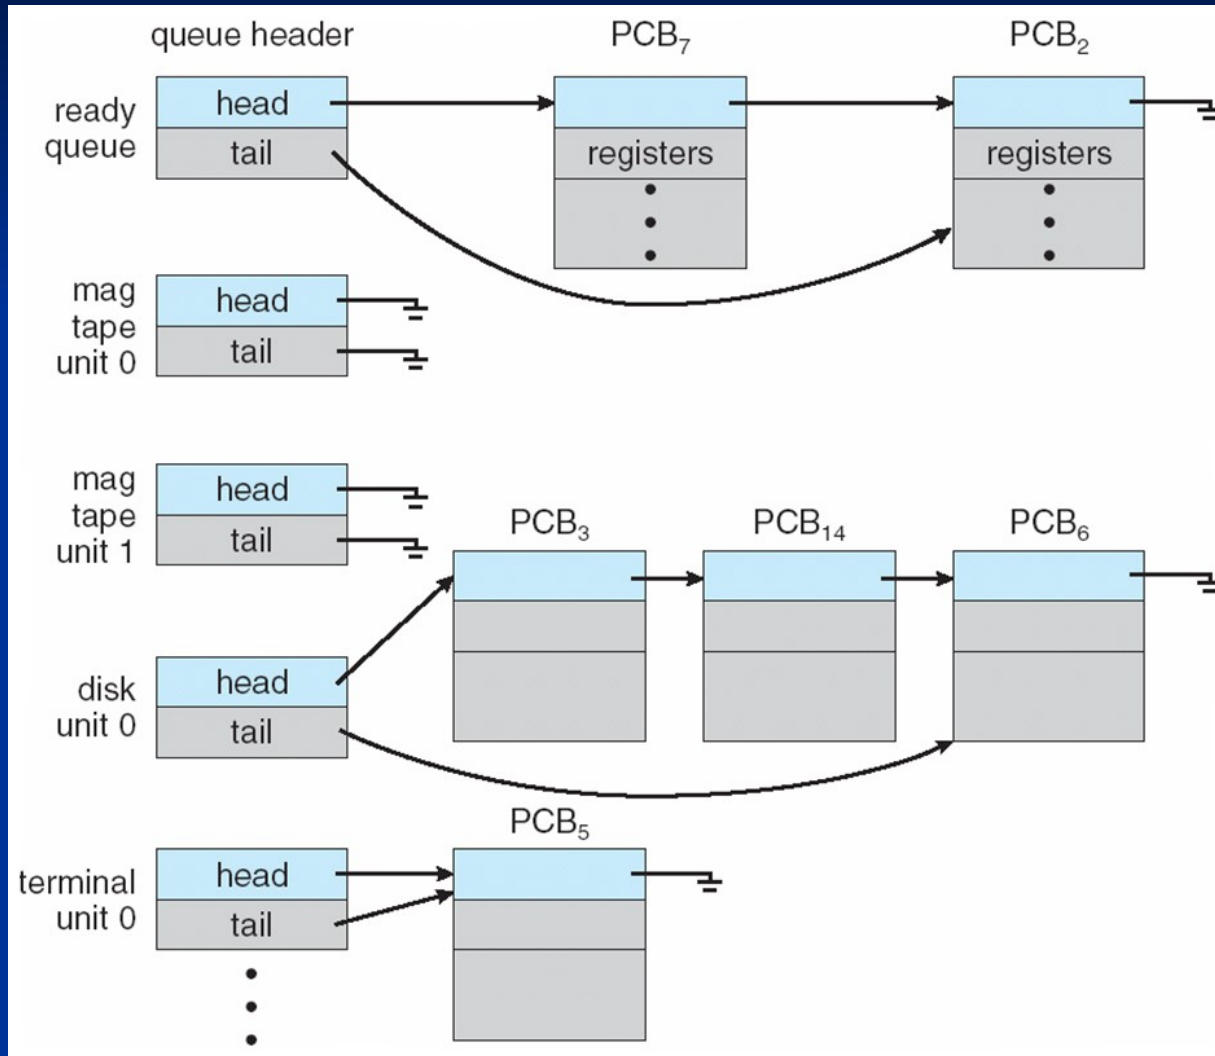
Threading

- Userland threads
 - POSIX Pthreads (IEEE 1003.1c)
 - Mach C-threads
 - Solaris threads
 - Windows threads
 - Java threads
- Kernel threads
 - Solaris LWP
 - Linux tasks (clone())
 - etc

Context switches

- TCB – Thread Control Block
 - Shared – parent process, execution time, memory, I/O resources
 - Private – PC, registers, stack, state information, pending/blocked signals
- TCB can be managed almost entirely in userland
 - Lower context switch overhead
- ...or may rely on the kernel in some way

Process queues



Process scheduling

- From user standpoint, OS permits many processes executing **simultaneously**
- In reality, OS switches among processes rapidly to give the **illusion** of simultaneity

Scheduler

- Operating System subsystem that is responsible for determining which process(es) to run, for how long, and when
- Two types: non-preemptive and preemptive

Non-preemptive

- Context switches happen **only** when the running process **waits or yields**
- Also called **cooperative multitasking**
- Used in Windows 3.1 and initial versions of MacOS

Preemptive

- Context switches can be forced
 - Usually after a fixed period of time, called a **quantum**
 - E.g., every 1/100sec
- Rely a **timer interrupt** that invokes the OS scheduler
 - Often the process that has been in the ready state the longest will execute next
- Implemented in *NIX, Windows 95 and above, etc

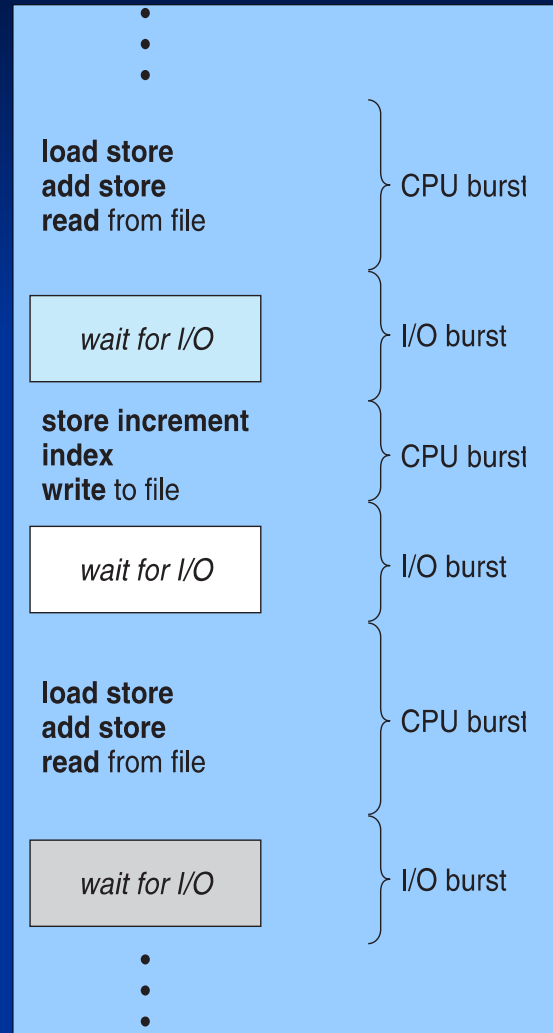
Tradeoffs

- Non-preemptive
 - More (user) control over how the CPU is used
 - Simple
- Preemptive
 - More robust
 - Enforced fairness


Preemptive scheduling

- Many different algorithms
 - FIFO, fixed priority, round-robin, multilevel queue, shortest remaining time first, worst case execution time, virtual round robin, etc

Execution cycle



CPU burst times




- 90% of CPU bursts are less than 10ms

First come first served (FCFS)

- Also known as FIFO

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Processes arrive in order (P_1, P_2, P_3)




- Waiting times: $P_1 = 0, P_2 = 24, P_3 = 27$
- Average wait time: $(0 + 24 + 27) / 3 = 17$

FCFS

- Suppose processes arrive:

P_2, P_3, P_1



- Waiting times: $P_1 = 6$, $P_2 = 0$, $P_3 = 3$
- Average waiting time:
$$(6 + 0 + 3) / 3 = 3$$
- **Convoy effect** – short process behind long process
 - Consider one CPU-bound and many I/O-bound processes

Purdue Trivia

- Purdue's "Big Bass Drum" (or BBD) was commissioned in 1921 by Paul Spotts Emrick
 - Leedy Manufacturing Company
- Ford Model T back axle and wheelbase

Round robin

- Each process gets a small unit of CPU time (a **quantum**)
 - Usually 10-100 milliseconds
- Quantum expires → preempt and move to end of the ready queue
 - Switch to the next process


Performance

- n processes in the ready queue
- time quantum q
- Each process will get $1/n$ of the CPU time in chunks of at most q time units at once
- \rightarrow Maximum wait time = $(n - 1) * q$
- What happens when q is large?
- What happens when it is small?

Round robin example

- Quantum = 4 ms

<u>Process</u>	<u>Burst Time</u>
P ₁	24
P ₂	3
P ₃	3



- Average turn around time:
 $(30 + 7 + 10) / 3 = 15.7 \text{ ms}$

Round robin example

- Quantum = 1 ms

<u>Process</u>	<u>Burst Time</u>
P ₁	24
P ₂	3
P ₃	3

- Average turn around time:
 $(8 + 9 + 29) / 3 = 15.3 \text{ ms}$

Quantum leap

- Shorter time quanta appear to reduce the average completion time
- What is the catch?
- Context switch overhead!

$$\% \text{overhead} = 100 * c / q$$

c = context switch overhead


- Suppose $c = 0.1\text{ms}$

$$q = 10 \rightarrow 1\%$$


$$q = 2 \rightarrow 5\%$$

$$q = 0.2 \rightarrow 50\%$$

Quantum and context switches



Turnaround time



Choosing a quantum

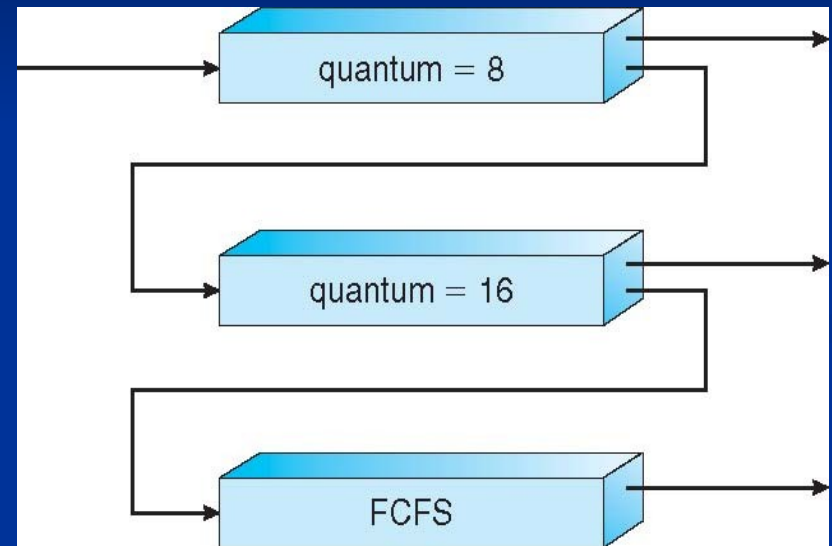
- Small enough to have a reasonable response time
- Large enough to control context switch overhead (should be $< 1\%$)
 - Shoot for at least 80% bursts $< q$
- 10ms is a good choice
 - 90% of bursts complete without preemption

Multilevel feedback queue

- Multiple queues, processes can move between them
- **Aging**: processes are moved to higher queue (increased priority) the longer they remain queued
- Reward interactive processes
 - Increase priority when a process yields before quantum expires
- Penalize CPU-bound processes

Example

- Suppose three queues:
 - Q_0 : RR, $q = 8\text{ms}$
 - Q_1 : RR, $q = 16\text{ ms}$
 - Q_2 : FCFS
- Scheduling
 - Process enters Q_0
 - Runs $> 8\text{ms}$, moves to Q_1
 - In Q_1 runs $> 16\text{ ms}$, moves to Q_2



- Reverse can happen
 - Process in Q_2 completes quickly (or ages) elevated to Q_1
 - Completes in $< 16\text{ms}$, elevates to Q_1

Questions?