



CS 25200: Systems Programming

Lecture 20: Threads and Synchronization

Prof. Turkstra



Lecture 20

- Threads
- Too much milk
- Mutual exclusion
- Semaphores

Creating threads

- POSIX

`pthread_create(&thr_id, attr, func, arg);`

- Solaris

`thr_create(stack, stack_size, func, arg, flags, &thr_id)`

- Windows

`CreateThread(attr, stack_size, func, arg, flags, &thr_id)`

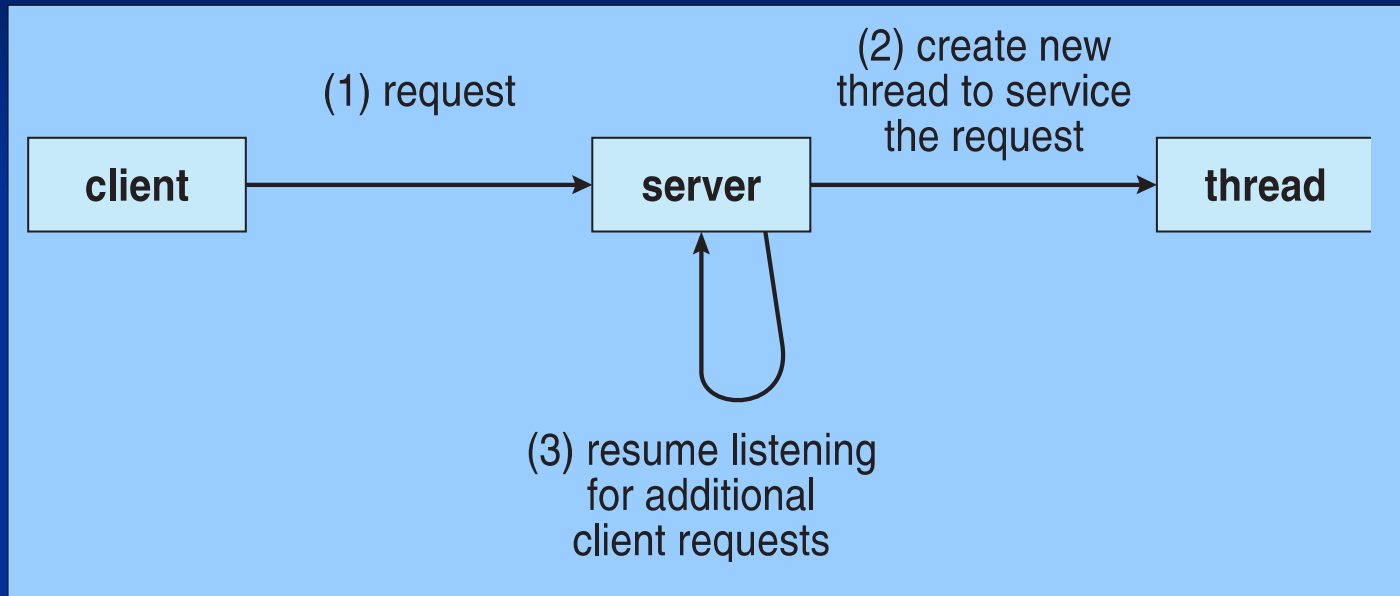
Remember

- Every thread has its own...
 - PC – program counter
 - Registers
 - State
 - Stack
- Process table entry manages this information for each thread

Why?

- Responsiveness – can continue executing if part of the process is blocked
 - Especially important for UIs
- Resource sharing – threads share resources automatically
 - Easier than explicit shared memory and message passing
- Economy – cheaper than process creation
 - Context switch overhead also lower
- Scalability – can leverage multiprocessor architectures

Concurrent server




- Suppose server is single threaded
 - What happens with slow clients (e.g., modems)?

Synchronization

- Instructions are guaranteed to execute in order within a single thread
 - Or at least appear to
- No ordering guaranteed among multiple threads
 - Threads may be interrupted at any time
 - Scheduling (“time slicing”)
- So?

Context switch



Too much milk

Terminology

- **Critical section:** section of code or collection of operations in which only one thread may be executing at a time
- **Mutual exclusion:** the property that exactly one thread is doing a certain thing at one time

What do we need?

- A locking mechanism
- Lock before using
- Unlock after done
- Wait if locked

More milk

Count example

```
#include <stdio.h>
#include <pthread.h>
```

```
int count;
```

```
void inc(long n) {
    for (int i = 0; i < n; i++) {
        count++;
    }
}
```

```
void dec(long n) {
    for (int i = 0; i < n; i++) {
        count--;
    }
}
```

Count example

```
int main(int argc, char **argv) {  
    long n = 10000000;  
    pthread_t t1;  
    pthread_t t2;  
  
    pthread_create(&t1, NULL, (void * (*)(void *)) inc, (void *) n);  
    pthread_create(&t2, NULL, (void * (*)(void *)) dec, (void *) n);  
  
    pthread_join(t1, NULL);  
    pthread_join(t2, NULL);  
  
    printf("%d\n", count);  
    return 0;  
}
```

Race condition

- `count++/--` statement is not atomic
 - load C
 - add one to C
 - store C
- Interleaving of thread statements impacts the result
 - Non-deterministic

Non-determinism

- Instruction execution order among separate threads depends on many things
 - Threading implementation
 - Operating system (scheduling, preempting)
 - Hardware interrupts
- This occurs even on single CPU, single core systems

Synchronization

- It is the programmer's job to anticipate all possible orderings and protect against errors
- Concurrency and synchronization is **hard**
 - Sometimes the overhead of implementing and debugging a concurrent program is not worth it
 - Therac-25
 - Safety critical systems should always have hardware interlocks

Purdue Trivia

- Orville Redenbacher graduated from Purdue in 1928 with a degree in Agronomy
 - Marched Tuba in the AAMB
 - Also on the track team
 - Worked for the exponent
 - Honorary doctorate in 1988
- Mural in PMU basement includes him



Criteria

■ Musts

- Processes **not** in critical section should not block others
- No one waits forever
- Multi-processor friendly

■ Desirable

- Fairness – everyone eventually gets into the critical section
- Efficient – don't waste resources (no busy waiting)
- Simple – symmetric code, easy to use
 - Like bracketing

Processes and mutual exclusion

- Always lock before manipulating shared memory
- Always unlock after manipulating shared memory
- Do not lock again if already locked
- Do not unlock if not locked by you
- Do not spend large amounts of time in critical section

Atomic

- Appears to the entire system as occurring all at once without interruption
 - No interrupts
 - No signals
 - No concurrent processes or threads

Semaphore

- Synchronization variable that takes on positive integer values
 - Dijkstra, 60s
- Two operations:
 - P(semaphore): atomic operation waits for semaphore > 0 , then **decrements** by one
 - “Proberen” in Dutch
 - V(semaphore): atomic operation **increments** by one
 - “Verhogen”

Pseudo code

P(S)

```
wait(S) {  
    while (S <= 0);  
    S--;  
}
```

V(S)

```
signal(S) {  
    S++;  
}
```

- Done atomically
- Not usually in hardware – implementation later

Binary semaphores

- Often called a mutex or lock
- Semaphore that takes on values of 0 and 1 only
- Too much milk?

```
P(milkSemaphore)
if (!milk)
    buy milk;
V(milkSemaphore)
```


Properties

- Machine-independent
- Simple
- Works with many processes
- Can have different critical sections with different semaphores
- Can acquire many resources simultaneously (multiple P's)
- Can permit multiple processes to enter critical section at once

Usage

- Mutual exclusion
 - One process is accessing critical section at a time
 - What about separate groups of data that need to be accessed independently
- Condition synchronization
 - Permit processes to wait for something
 - What if disparate groups of processes want to wait for unrelated events?

Fixing inc.c

Implementation

- Uniprocessor solution: disable interrupts!

```
typedef struct {  
    int count;  
    queue q;  
} semaphore;
```

P

```
void P(semaphore S) {  
    disable interrupts;  
    if (s->count > 0) {  
        s->count--;  
        enable interrupts;  
        return;  
    }  
    add(s->q, current_thread);  
    sleep(); // re-dispatch  
    enable interrupts;  
    return;  
}
```

V

```
void V(semaphore S) {  
    disable interrupts;  
    if (isEmpty(s->q)) {  
        s->count++;  
    } else {  
        thread = removeFirst(s->q);  
        wakeup(thread); // put thread on ready q  
    }  
    enable interrupts;  
    return;  
}
```

Multiprocessor?

- Cannot just turn off interrupts
 - Doesn't prevent other processors from accessing shared memory
- Turn off other processors?
 - Bad :-(
- Use atomic read and write?
 - Needs to be atomic across all processors!
- Big research area for a long time

Test-and-set (IBM)

- Atomic read-modify-write instruction
- TAS – on most CISC architectures
- Semantics:
 - Set value to k, but return old value
 - $k = 1 \rightarrow$ binary semaphore

```
int lock;  
while (TAS(&lock, 1) != 0);  
<critical section>  
lock = 0;
```


TAS

- Implemented by memory hardware or CPU refusing to relinquish bus access
- Still have to disable interrupts on current core
- Why?

RISC Mechanism

- Load-linked
 - ldl – loads a word from memory and sets per-processor flag associated with that word (in cache)
 - Store operation to same location (by any processor) resets all processors' flags for that word
- Store-conditionally
 - stc – stores word iff flag still set, indicates success or failure

```
int lock;  
while ((ldl(&lock) != 0) || !stc(&lock, 1));  
<critical section>  
lock = 0;
```

Questions?