



CS 25200: Systems Programming

Lecture 21: Semaphore Implementation, Thread Safety and Race Conditions

Prof. Turkstra



Midterm Exam

- Thursday, October 17 8:00pm – 10:00pm
 - WALC 1055
 - Covers material through and including lecture 19 (scheduling)
 - Lectures and labs
 - Homework out tonight/tomorrow
 - Due before exam
 - Print and hand in
 - Not a comprehensive indicator

Announcements

- Lab 3 Checkpoint grades available
- Quiz/attendance too
- Remember: one week regrade request window!

Lecture 21

- Semaphore review
- Semaphore implementation
- Spinlocks
- Threads and `fork()`
- Thread safety
- Races

Terminology

- **Critical section:** section of code or collection of operations in which only one thread may be executing at a time
- **Mutual exclusion:** the property that exactly one thread is doing a certain thing at one time

Count example

```
#include <stdio.h>
#include <pthread.h>
```

```
int count;
```

```
void inc(long n) {
    for (int i = 0; i < n; i++) {
        count++;
    }
}
```

```
void dec(long n) {
    for (int i = 0; i < n; i++) {
        count--;
    }
}
```

Count example

```
int main(int argc, char **argv) {  
    long n = 10000000;  
    pthread_t t1;  
    pthread_t t2;  
  
    pthread_create(&t1, NULL, (void * (*)(void *)) inc, (void *) n);  
    pthread_create(&t2, NULL, (void * (*)(void *)) dec, (void *) n);  
  
    pthread_join(t1, NULL);  
    pthread_join(t2, NULL);  
  
    printf("%d\n", count);  
    return 0;  
}
```

Race condition

- `count++/--` statement is not atomic
 - load C
 - add one to C
 - store C
- Interleaving of thread statements impacts the result
 - Non-deterministic

Non-determinism

- Instruction execution order among separate threads depends on many things
 - Threading implementation
 - Operating system (scheduling, preempting)
 - Hardware interrupts
- This occurs even on single CPU, single core systems

Synchronization

- It is the programmer's job to anticipate all possible orderings and protect against errors
- Concurrency and synchronization is **hard**
 - Sometimes the overhead of implementing and debugging a concurrent program is not worth it
 - Therac-25
 - Safety critical systems should always have hardware interlocks

Criteria

■ Musts

- Processes **not** in critical section should not block others
- No one waits forever
- Multi-processor friendly

■ Desirable

- Fairness – everyone eventually gets into the critical section
- Efficient – don't waste resources (no busy waiting)
- Simple – symmetric code, easy to use
 - Like bracketing

Processes and mutual exclusion

- Always lock before manipulating shared memory
- Always unlock after manipulating shared memory
- Do not lock again if already locked
- Do not unlock if not locked by you
- Do not spend large amounts of time in critical section

Atomic

- Appears to the entire system as occurring all at once without interruption
 - No interrupts
 - No signals
 - No concurrent processes or threads

Semaphore

- Synchronization variable that takes on positive integer values
 - Dijkstra, 60s
- Two operations:
 - P(semaphore): atomic operation waits for semaphore > 0 , then **decrements** by one
 - “Proberen” in Dutch
 - V(semaphore): atomic operation **increments** by one
 - “Verhogen”

Pseudo code

P(S)

```
wait(S) {  
    while (S <= 0);  
    S--;  
}
```

V(S)

```
signal(S) {  
    S++;  
}
```

- Done atomically
- Not usually in hardware – implementation later

Binary semaphores

- Often called a mutex or lock
- Semaphore that takes on values of 0 and 1 only
- Too much milk?

```
P(milkSemaphore)
if (!milk)
    buy milk;
V(milkSemaphore)
```


Properties

- Machine-independent
- Simple
- Works with many processes
- Can have different critical sections with different semaphores
- Can acquire many resources simultaneously (multiple P's)
- Can permit multiple processes to enter critical section at once

Usage

- Mutual exclusion
 - One process is accessing critical section at a time
 - What about separate groups of data that need to be accessed independently
- Condition synchronization
 - Permit processes to wait for something
 - What if disparate groups of processes want to wait for unrelated events?

Fixing inc.c

Implementation

- Uniprocessor solution: disable interrupts!

```
typedef struct {  
    int count;  
    queue q;  
} semaphore;
```

P

```
void P(semaphore s) {  
    disable interrupts;  
    if (s->count > 0) {  
        s->count--;  
        enable interrupts;  
        return;  
    }  
    add(s->q, current_thread);  
    sleep(); // re-dispatch  
    enable interrupts;  
    return;  
}
```

V

```
void V(semaphore s) {  
    disable interrupts;  
    if (isEmpty(s->q)) {  
        s->count++;  
    } else {  
        thread = removeFirst(s->q);  
        wakeup(thread); // put thread on ready q  
    }  
    enable interrupts;  
    return;  
}
```

Multiprocessor?

- Cannot just turn off interrupts
 - Doesn't prevent other processors from accessing shared memory
- Turn off other processors?
 - Bad :-(
- Use atomic read and write?
 - Needs to be atomic across all processors!
- Big research area for a long time

Test-and-set (IBM)

- Atomic read-modify-write instruction
- TAS – on most CISC architectures
- Semantics:
 - Set value to k, but return old value
 - $k = 1 \rightarrow$ binary semaphore

```
int lock;  
while (TAS(&lock, 1) != 0);  
<critical section>  
lock = 0;
```


TAS

- Implemented by memory hardware or CPU refusing to relinquish bus access
- Still have to disable interrupts on current core
- Why?

RISC Mechanism

- Load-linked
 - ldl – loads a word from memory and sets per-processor flag associated with that word (in cache)
 - Store operation to same location (by any processor) resets all processors' flags for that word
- Store-conditionally
 - stc – stores word iff flag still set, indicates success or failure

```
int lock;  
while ((ldl(&lock) != 0) || !stc(&lock, 1));  
<critical section>  
lock = 0;
```

Purdue Trivia

- We are on the seventh iteration of the Boilermaker Special, Purdue's official mascot.
 - World's largest, fastest, heaviest, and loudest collegiate mascot
 - Dedicated 9/3/2011
- Boilermaker Xtra Special is a smaller version designed for use indoors
 - Eighth iteration
- Both are entrusted to the Purdue Reamer Club



Spinlocks

- The two previous solutions rely on a **spinlock**
 - Busy waiting
- Still used for multicore implementations
 - Optimization: if thread holding lock is sleeping, also sleep

Mutexes

- Spinlocks are not guaranteed to be fair
- Can build semaphores on top of spinlocks
 - Similar to single-processor implementation
 - Replace enable/disable interrupts with `spin_lock()` and `spin_unlock()`

Threads and fork()

- `fork()` behaves differently on some platforms
- For most *nix variants, calling `fork()` duplicates only the calling thread
 - Even if that thread is the parent of other threads
- Solaris' `thr_create()` will also duplicate the children
 - `fork1()` if you wish to avoid that

Thread safety

- Thread-safe code manipulates shared data structures properly
- Thread safe: guaranteed to be free of **race conditions** when accessed or used by multiple threads **simultaneously**
- Not thread safe: should **not** be used simultaneously by multiple threads
- Often designated in man pages

man page

- Function may be marked MT-Unsafe
 - strtok()
- Or MT-Safe
 - strtok_r()
- Often times MT-Safe code is **re-entrant** – state is saved e.g. on a stack

Linked list race

```
typedef struct linked_list {  
    struct node *head;  
} linked_list;
```

```
typedef struct node {  
    int value;  
    struct node *next;  
} node;
```

```
linked_list *list = malloc(sizeof(list));
```

```
int insert(linked_list *list, int val) {  
    if (!list->head) {  
        list->head = malloc(sizeof(node));  
        list->head->value = val;  
        return;  
    }  
  
    node *new_node = malloc(sizeof(node));  
    new_node->next = list->head;  
    list->head = new_node;  
}
```

Race condition

- Output is dependent on the sequence or timing of uncontrollable events

Thread safe

```
typedef struct linked_list {  
    pthread_mutex_t mutex;  
    struct node *head;  
} linked_list;
```

```
typedef struct node {  
    int value;  
    struct node *next;  
} node;
```

```
linked_list *list = malloc(sizeof(linked_list));  
pthread_mutex_init(&list->mutex, NULL);
```

```
int insert(linked_list *list, int val) {  
    pthread_mutex_lock(&list->mutex);  
    if (!list->head) {  
        list->head = malloc(sizeof(node));  
        list->head->value = val;  
        pthread_mutex_unlock(&list->mutex);  
        return;  
    }  
}
```

```
node *new_node = malloc(sizeof(node));  
new_node->next = list->head;  
list->head = new_node;  
pthread_mutex_unlock(&list->mutex);  
}
```

Remove

```
int remove_head(linked_list *list) {  
    node *tmp = NULL;  
    int val;  
  
    tmp = list->head;  
    if (tmp == NULL) return -1;  
    list->head = list->head->next;  
    val = tmp->value;  
    free(tmp);  
    tmp = NULL;  
  
    return val;  
}
```

Remove MT-Safe

```
int remove_head(linked_list *list) {
    node *tmp = NULL;
    int val;

    pthread_mutex_lock(&list->mutex);
    tmp = list->head;
    if (tmp == NULL) {
        pthread_mutex_unlock(&list->mutex);
        return -1;
    }
    list->head = list->head->next;
    pthread_mutex_unlock(&list->mutex);
    val = tmp->value;
    free(tmp);
    tmp = NULL;

    return val;
}
```


What if?

- What if we forgot the first call to `pthread_mutex_unlock()`?
- What if we move `pthread_mutex_unlock()` to right above the return statement?
- What if we had a mutex for every node?



Suggested homework

- Analyze singly linked list, doubly linked list, and binary tree algorithms
 - Where are the possible synchronization bugs
 - How can you prevent them?
 - What might the trade-offs be?

Questions?